


CS 5522: Artificial Intelligence II

Final Review



Instructor: Wei Xu

Ohio State University

[Cartoons were adapted from CS188 Intro to AI at UC Berkeley.]

Final Exam

- The final exam will be closed notes, books, laptops, and people. (You will be given an instructor-provided “cheat sheet”)
- 105 minutes. 12/12 Wednesday 12:00-1:45pm. This room.
- Preparation:
 - Lecture Slides
 - Written Homework
 - Example Exam
 - Project 3: Ghostbusters


Final Exam

- Make sure you **understand** the fundamentals in addition to being able to procedurally execute algorithms.
- The exam will not test your knowledge of Python, however questions may assume **familiarity with the projects** and test ability of **writing pseudocode**.
- See **written homework** and **example exam** for examples

Possible Final Exam Topics

- Markov Models and HMMs:

- Markov Model and HMM definition
- Observation and Time Elapse
- Exact Inference
- Approximate Inference
- Particle Filtering
- Most Likely Explanation




- Bayes Nets:

- Probabilistic Representations
- D-Separation
- Inference
- Variable Elimination

- Naive Bayes:

- Classification
- Prediction
- Parameter Estimation
- Smoothing



Possible Final Exam Topics


- Will not cover the follows:
 - Search
 - Reinforcement Learning
 - Neural Networks
 - Speech Recognition
 - Computer Vision
 - Natural Language Processing

Office Hours

- Office Hour - Wednesday, 12/5 4-5pm (DL 495)
- Office Hour - Friday, 12/7 4-5pm (DL 495)
- Will also answer questions on Piazza

Instructor provided “cheat sheet”


Inference: Base Cases



Passage of Time

- Assume we have current belief $P(X \mid \text{evidence to date})$

$$B(X_t) = P(X_t | e_{1:t})$$



- Then, after one time step passes:

$$\begin{aligned} P(X_{t+1} | e_{1:t}) &= \sum_{x_t} P(X_{t+1}, x_t | e_{1:t}) \\ &= \sum_{x_t} P(X_{t+1} | x_t, e_{1:t}) P(x_t | e_{1:t}) \\ &= \sum_{x_t} P(X_{t+1} | x_t) P(x_t | e_{1:t}) \end{aligned}$$

- Or compactly:

$$B'(X_{t+1}) = \sum_{x_t} P(X' | x_t) B(x_t)$$


- Basic idea: beliefs get “pushed” through the transitions

- With the “B” notation, we have to be careful about what time step t the belief is about, and what evidence it includes

Active / Inactive Paths

- Question: Are X and Y conditionally independent given evidence variables $\{Z\}$?

- Yes, if X and Y “d-separated” by Z
- Consider all (undirected) paths from X to Y
- No active paths = independence!



- A path is active if each triple is active:

- Causal chain $A \rightarrow B \rightarrow C$ where B is unobserved (either direction)
- Common cause $A \leftarrow B \rightarrow C$ where B is unobserved
- Common effect (aka v-structure) $A \rightarrow B \leftarrow C$ where B or one of its descendants is observed

- All it takes to block a path is a single inactive segment

Observation

- Assume we have current belief $P(X \mid \text{previous evidence})$:

$$B'(X_{t+1}) = P(X_{t+1} | e_{1:t})$$


- Then, after evidence comes in:

$$\begin{aligned} P(X_{t+1} | e_{1:t+1}) &= P(X_{t+1}, e_{t+1} | e_{1:t}) / P(e_{t+1} | e_{1:t}) \\ &\propto_{X_{t+1}} P(X_{t+1}, e_{t+1} | e_{1:t}) \\ &= P(e_{t+1} | e_{1:t}, X_{t+1}) P(X_{t+1} | e_{1:t}) \\ &= P(e_{t+1} | X_{t+1}) P(X_{t+1} | e_{1:t}) \end{aligned}$$

- Or, compactly:

$$B(X_{t+1}) \propto_{X_{t+1}} P(e_{t+1} | X_{t+1}) B'(X_{t+1})$$

- Basic idea: beliefs “reweighted” by likelihood of evidence
- Unlike passage of time, we have to renormalize



The Forward Algorithm

- We are given evidence at each time and want to know

$$B_t(X) = P(X_t | e_{1:t})$$

- We can derive the following updates

$$\begin{aligned} P(x_t | e_{1:t}) &\propto_X P(x_t, e_{1:t}) \\ &= \sum_{x_{t-1}} P(x_{t-1}, x_t, e_{1:t}) \\ &= \sum_{x_{t-1}} P(x_{t-1}, e_{1:t-1}) P(x_t | x_{t-1}) P(e_t | x_t) \\ &= P(e_t | x_t) \sum_{x_{t-1}} P(x_{t-1} | e_{1:t-1}) P(x_{t-1}, e_{1:t-1}) \end{aligned}$$

We can normalize as we go if we want to have $P(x|e)$ at each time step, or just once at the end...

- Forward Algorithm
- D-Separation