


CNNs and Neural CRFs

Wei Xu


(many slides from Greg Durrett)

Recall: RNNs

- ▶ Cell that takes some input x , has some hidden state h , and updates that hidden state and produces output y (all vector-valued)



Recall: RNN Abstraction



the movie was great

- ▶ **Encoding of the sentence** – can pass this a decoder or make a classification decision about the sentence
- ▶ **Encoding of each word** – can pass this to another layer to make a prediction (can also pool these to get a different sentence encoding)
- ▶ RNN can be viewed as a transformation of a sequence of vectors into a sequence of context-dependent vectors

What can LSTMs model?

- ▶ Sentiment
 - ▶ Encode one sentence, predict
- ▶ Language models
 - ▶ Move left-to-right, per-token prediction
- ▶ Translation
 - ▶ Encode sentence + then decode, use token predictions for attention weights (next lecture)
- ▶ Textual entailment
 - ▶ Encode two sentences, predict

Natural Language Inference

Premise		Hypothesis
A boy plays in the snow	<i>entails</i>	A boy is outside
A man inspects the uniform of a figure	<i>contradicts</i>	The man is sleeping
An older and younger man smiling	<i>neutral</i>	Two men are smiling and laughing at cats playing

- ▶ Long history of this task: “Recognizing Textual Entailment” challenge in 2006 (Dagan, Glickman, Magnini)
- ▶ Early datasets: small (hundreds of pairs), very ambitious (lots of world knowledge, temporal reasoning, etc.)

SNLI Dataset

- ▶ Show people captions for (unseen) images and solicit entailed / neural / contradictory statements

- ▶ >500,000 sentence pairs

- ▶ Encode each sentence and process

100D LSTM: 78% accuracy


300D LSTM: 80% accuracy

(Bowman et al., 2016)

300D BiLSTM: 83% accuracy

(Liu et al., 2016)

- ▶ Later: better models for this



Bowman et al. (2015)

This Lecture


- ▶ CNNs
- ▶ CNNs for Sentiment
- ▶ Neural CRFs

CNNs


Convolutional Layer

- ▶ Applies a *filter* over patches of the input and returns that filter's activations
- ▶ Convolution: take dot product of filter with a patch of the input

image: $n \times n \times k$



filter: $m \times m \times k$



sum over dot products

$$\text{activation}_{ij} = \sum_{i_o=0}^{k-1} \sum_{j_o=0}^{k-1} \text{image}(i + i_o, j + j_o) \cdot \text{filter}(i_o, j_o)$$

↑
offsets


Each of these cells is a vector with multiple values

Images: RGB values (3 dim)

Convolutional Layer

- ▶ Applies a *filter* over patches of the input and returns that filter's activations
- ▶ Convolution: take dot product of filter with a patch of the input

image: $n \times n \times k$ filter: $m \times m \times k$ activations: $(n - m + 1) \times (n - m + 1) \times 1$



Convolutions for NLP


- ▶ Input and filter are 2-dimensional instead of 3-dimensional

sentence: n words $\times k$ vec dim

filter: $m \times k$


activations: $(n - m + 1) \times 1$

the movie was good



- ▶ Combines evidence locally in a sentence and produces a new (but still variable-length) representation


Compare: CNNs vs. LSTMs



$O(n) \times c$




c filters,
 $m \times k$ each



$n \times k$

the movie was good



$n \times 2c$

BiLSTM with
hidden size c


$n \times k$

the movie was good

- ▶ Both LSTMs and convolutional layers transform the input using context
- ▶ LSTM: “globally” looks at the entire sentence (but local for many problems)
- ▶ CNN: local depending on filter width + number of layers

CNNs for Sentiment

CNNs for Sentiment Analysis



$$P(y|x)$$

projection + softmax

c -dimensional vector

max pooling over the sentence

$$n \times c$$


c filters,
 $m \times k$ each

$$n \times k$$

the movie was good


- ▶ Max pooling: return the max activation of a given filter over the entire sentence; like a logical OR (sum pooling is like logical AND)

Understanding CNNs for Sentiment




- ▶ Filter “looks like” the things that will cause it to have high activation

Understanding CNNs for Sentiment




Understanding CNNs for Sentiment




- ▶ Takes variable-length input and turns it into fixed-length output
- ▶ Filters are initialized randomly and then learned

Understanding CNNs for Sentiment



- ▶ Word vectors for similar words are similar, so convolutional filters will have similar outputs

Understanding CNNs for Sentiment



- ▶ Analogous to bigram features in bag-of-words models
- ▶ Indicator feature of text containing bigram \leftrightarrow max pooling of a filter that matches that bigram

What can CNNs learn?

the movie was not good



the movie was not really all that good

the cinematography was good, the music great, but the movie was bad

I entered the theater in the bloom of youth and left as an old man

Deep Convolutional Networks



- ▶ Low-level filters: extract low-level features from the data



Zeiler and Fergus (2014)

Deep Convolutional Networks

- ▶ High-level filters: match larger and more “semantic patterns”




Zeiler and Fergus (2014)

CNNs: Implementation

- ▶ Input is $\text{batch_size} \times n \times k$ matrix, filters are $c \times m \times k$ matrix (c filters)
- ▶ Typically use filters with m ranging from 1 to 5 or so (multiple filter widths in a single convnet)
- ▶ All computation graph libraries support efficient convolution operations

CNNs for Sentence Classification

- ▶ Question classification, sentiment, etc.
- ▶ Conv+pool, then use feedforward layers to classify
- ▶ Can use multiple types of input vectors (fixed initializer and learned)



the movie was good

Kim (2014)

Sentence Classification

Model	MR	SST-1	SST-2	Subj	TREC	CR	MPQA
CNN-multichannel	81.1	47.4	88.1	93.2	92.2	85.0	89.4
NBSVM (Wang and Manning, 2012)	79.4	-	-	93.2	-	81.8	86.3

movie review
sentiment

subjectivity/objectivity
detection

product
reviews

question type
classification

```
graph TD; A[Movie review sentiment] --> B[Model]; C[Subjectivity/objectivity detection] --> B; D[Product reviews] --> B; E[Question type classification] --> B; B --- R1[NBSVM]; B --- R2[CNN];
```

- ▶ Also effective at document-level text classification

Neural CRF Basics

NER Revisited

- ▶ Features in CRFs: $I[\text{tag}=\text{B-LOC} \ \& \ \text{curr_word}=\text{Hangzhou}]$,
 $I[\text{tag}=\text{B-LOC} \ \& \ \text{prev_word}=to]$, $I[\text{tag}=\text{B-LOC} \ \& \ \text{curr_prefix}=\text{Han}]$
 - ▶ Linear model over features
 - ▶ Downsides:
 - ▶ Lexical features mean that words need to be seen in the training data
 - ▶ Linear model can't capture feature conjunctions as effectively (doesn't work well to look at more than 2 words with a single feature)

LSTMs for NER

B-PER I-PER O O O B-LOC O O O B-ORG O O


Barack Obama will travel to Hangzhou today for the G20 meeting.

PERSON

LOC

ORG

B-PER I-PER O O O B-LOC



Barack Obama will travel to Hangzhou

- ▶ Transducer (LM-like model)
- ▶ What are the strengths and weaknesses of this model compared to CRFs?

LSTMs for NER

B-PER	I-PER	0	0	0	B-LOC	0	0	0	B-ORG	0	0
-------	-------	---	---	---	-------	---	---	---	-------	---	---


Barack Obama will travel to Hangzhou today for the G20 meeting .

PERSON

LOC

ORG

B-PER	I-PER	0	0	0	B-LOC
-------	-------	---	---	---	-------



Barack Obama will travel to Hangzhou

- ▶ Bidirectional transducer model
- ▶ What are the strengths and weaknesses of this model compared to CRFs?

Neural CRFs


B-PER I-PER 0 0 0 B-LOC 0 0 0 B-ORG 0 0

Barack Obama will travel to Hangzhou today for the G20 meeting.

PERSON

LOC

ORG




Barack Obama will travel to Hangzhou

- ▶ Neural CRFs: bidirectional LSTMs (or some NN) compute emission potentials, capture structural constraints in transition potentials

Neural CRFs


$$P(\mathbf{y}|\mathbf{x}) = \frac{1}{Z} \prod_{i=2}^n \exp(\phi_t(y_{i-1}, y_i)) \prod_{i=1}^n \exp(\phi_e(y_i, i, \mathbf{x}))$$



- ▶ Conventional: $\phi_e(y_i, i, \mathbf{x}) = w^\top f_e(y_i, i, \mathbf{x})$
- ▶ Neural: $\phi_e(y_i, i, \mathbf{x}) = W_{y_i}^\top f(i, \mathbf{x})$ W is a `num_tags x len(f)` matrix
- ▶ $f(i, \mathbf{x})$ could be the output of a feedforward neural network looking at the words around position i , or the i th output of an LSTM, ...
- ▶ Neural network computes unnormalized potentials that are consumed and “normalized” by a structured model
- ▶ Inference: compute f , use Viterbi

Computing Gradients

$$P(\mathbf{y}|\mathbf{x}) = \frac{1}{Z} \prod_{i=2}^n \exp(\phi_t(y_{i-1}, y_i)) \prod_{i=1}^n \exp(\phi_e(y_i, i, \mathbf{x}))$$



- ▶ Conventional: $\phi_e(y_i, i, \mathbf{x}) = w^\top f_e(y_i, i, \mathbf{x})$
- ▶ Neural: $\phi_e(y_i, i, \mathbf{x}) = W_{y_i}^\top f(i, \mathbf{x})$
- ▶ For linear model: $\frac{\partial \mathcal{L}}{\partial \phi_{e,i}} = -P(y_i = s|\mathbf{x}) + I[s \text{ is gold}]$ “error signal”, compute with F-B
 - ▶ chain rule say to multiply together, gives our update
- ▶ For neural model: compute gradient of phi w.r.t. parameters of neural net

Neural CRFs


B-PER I-PER 0 0 0 B-LOC 0 0 0 B-ORG 0 0

Barack Obama will travel to Hangzhou today for the G20 meeting.

PERSON

LOC

ORG



Barack Obama will travel to Hangzhou

- 1) Compute $f(x)$
- 2) Run forward-backward
- 3) Compute error signal
- 4) Backprop (no knowledge of sequential structure required)

FFNN Neural CRF for NER


B-PER I-PER 0 0 0 B-LOC 0 0 0 B-ORG 0 0

Barack Obama will travel to Hangzhou today for the G20 meeting.

PERSON

LOC

ORG



$$\phi_e = Wg(Vf(\mathbf{x}, i))$$

$$f(\mathbf{x}, i) = [emb(\mathbf{x}_{i-1}), emb(\mathbf{x}_i), emb(\mathbf{x}_{i+1})]$$

previous word curr word next word

to Hangzhou today

LSTM Neural CRFs


B-PER	I-PER	0	0	0	B-LOC	0	0	0	B-ORG	0	0
-------	-------	---	---	---	-------	---	---	---	-------	---	---

Barack Obama will travel to Hangzhou today for the G20 meeting .

PERSON

LOC

ORG



Barack Obama will travel to Hangzhou

- ▶ Bidirectional LSTMs compute emission (or transition) potentials

LSTMs for NER

B-PER I-PER O O O B-LOC O O O B-ORG O O


Barack Obama will travel to Hangzhou today for the G20 meeting.

PERSON

LOC

ORG

B-PER I-PER O O O B-LOC




Barack Obama will travel to Hangzhou

- ▶ How does this compare to neural CRF?


“NLP (Almost) From Scratch”

Approach	POS (PWA)	CHUNK (F1)	NER (F1)	SRL (F1)
Benchmark Systems	97.24	94.29	89.31	77.92
NN+WLL	96.31	89.13	79.53	55.40
NN+SLL	96.37	90.33	81.47	70.99
NN+WLL+LM1	97.05	91.91	85.68	58.18
NN+SLL+LM1	97.10	93.65	87.58	73.84
NN+WLL+LM2	97.14	92.04	86.96	58.34
NN+SLL+LM2	97.20	93.63	88.67	74.15

- ▶ WLL: independent classification; SLL: neural CRF
- ▶ LM2: word vectors learned from a precursor to word2vec/GloVe, trained for 2 weeks (!) on Wikipedia



CNN Neural CRFs



- ▶ Append to each word vector an *embedding of the relative position* of that word
- ▶ Convolution over the sentence produces a position-dependent representation



CNN NCRFs vs. FFNN NCRFs

Approach	POS (PWA)	CHUNK (F1)	NER (F1)	SRL (F1)
Benchmark Systems	97.24	94.29	89.31	77.92
<i>Window Approach</i>				
NN+SLL+LM2	97.20	93.63	88.67	-
<i>Sentence Approach</i>				
NN+SLL+LM2	97.12	93.37	88.78	74.15

- ▶ Sentence approach (CNNs) is comparable to window approach (FFNNs) except for SRL where they claim it works much better

Neural CRFs with LSTMs

- ▶ Neural CRF using character LSTMs to compute word representations



Neural CRFs with LSTMs

- ▶ Chiu+Nichols: character CNNs instead of LSTMs
- ▶ Lin/Passos/Luo: use external resources like Wikipedia
- ▶ LSTM-CRF captures the important aspects of NER: word context (LSTM), sub-word features (character LSTMs), outside knowledge (word embeddings)

Model	F ₁
Collobert et al. (2011)*	89.59
Lin and Wu (2009)	83.78
Lin and Wu (2009)*	90.90
Huang et al. (2015)*	90.10
Passos et al. (2014)	90.05
Passos et al. (2014)*	90.90
Luo et al. (2015)* + gaz	89.9
Luo et al. (2015)* + gaz + linking	91.2
Chiu and Nichols (2015)	90.69
Chiu and Nichols (2015)*	90.77
LSTM-CRF (no char)	90.20
LSTM-CRF	90.94

Takeaways

- ▶ CNNs are a flexible way of extracting features analogous to bag of n-grams, can also encode positional information
- ▶ All kinds of NNs can be integrated into CRFs for structured inference. Can be applied to NER, other tagging, parsing, ...